

Objects in Linux

Trust us, we know what we are doing...

Greg Kroah-Hartman
gregkh@linuxfoundation.org



/THEORY/IN/PRACTICE

Beautiful Code

Leading Programmers Explain How They Think

O'REILLY®

Edited by Andy Oram & Greg Wilson

Early 2000's

“Unify all Linux devices”

Pat Mochel – OSDL

Greg K-H - IBM

```
struct device {
    struct list_head node;
    struct list_head children;
    struct device *parent;

    char name[DEVICE_NAME_SIZE];
    bus_id[BUS_ID_SIZE];

    spinlock_t lock;

    atomic_t refcount;

    struct driver_dir_entry *dir;
    struct device_driver *driver;

    void *driver_data;
    void *platform_data;
    u32 current_state;
    unsigned char *saved_state;
};
```

```
struct device {
    struct list_head node;
    struct list_head children;
    struct device *parent;

    char name[DEVICE_NAME_SIZE];
    bus_id[BUS_ID_SIZE];

    spinlock_t lock;

    atomic_t refcount;

    struct driver_dir_entry *dir;
    struct device_driver *driver;

    void *driver_data;
    void *platform_data;
    u32 current_state;
    unsigned char *saved_state;
};
```





github.com/gregkh/presentation-kref

Objects in Linux

Trust us, we know what we are doing...

Greg Kroah-Hartman
gregkh@linuxfoundation.org





Talk based on a chapter I wrote for
“Beautiful Code – Leading Programmers
Explain How They Think”

Edited by Andy Oram and Greg Wilson
O'Reilly
2007

Early 2000's “Unify all Linux devices”

Pat Mochel – OSDL
Greg K-H - IBM

Pat wanted this for power management and suspend/resume

I wanted this for persistent device naming.

All devices and subsystems were islands

Both tasks needed a way to see all devices in the system, suspend/resume wanted to know which device to suspend in which order.

Naming needed a way to assign a character or block device to a specific hardware device

```
struct device {
    struct list_head node;
    struct list_head children;
    struct device *parent;

    char name[DEVICE_NAME_SIZE];
    bus_id[BUS_ID_SIZE];

    spinlock_t lock;

    atomic_t refcount;

    struct driver_dir_entry *dir;
    struct device_driver *driver;

    void *driver_data;
    void *platform_data;
    u32 current_state;
    unsigned char *saved_state;
};
```

We came up with 'struct device'

All busses in the kernel were changed to create a structure based on this one. It was passed to the new driver core, and the driver core created a hierarchy of everything in the kernel.

This can be seen in sysfs (which used to be called driverfs)

```
struct device {
    struct list_head node;
    struct list_head children;
    struct device *parent;

    char name[DEVICE_NAME_SIZE];
    bus_id[BUS_ID_SIZE];

    spinlock_t lock;

    atomic_t refcount;

    struct driver_dir_entry *dir;
    struct device_driver *driver;

    void *driver_data;
    void *platform_data;
    u32 current_state;
    unsigned char *saved_state;
};
```

We came up with 'struct device'

All busses in the kernel were changed to create a structure based on this one. It was passed to the new driver core, and the driver core created a hierarchy of everything in the kernel.

This can be seen in sysfs (which used to be called driverfs)



github.com/gregkh/presentation-kref

Obligatory Penguin Picture

